

F RATIO UNDER THE NULL HYPOTHESIS

Packages used in this notebook

- `using Pkg`
- `Pkg.activate("/Users/juan/Documents/Julia/Tutorials/DataScience/Project.toml")`
- `using StatsPlots`
- `using Distributions`
- `using Random`
- `using DataFrames`
- `using StatsBase`
- `using Statistics`
- `using WebIO`
- `using PlutoUI`

Table of Contents

F RATIO UNDER THE NULL HYPOTHESIS

Packages used in this notebook

F distribution

Data

Summary statistics

F ratio

Generating a sampling distribution by reassignment under the null hypothesis

Estimating probability of our F ratio

- `PlutoUI.TableOfContents(title=" Table of Contents", indent=true, depth=4, aside=true)`

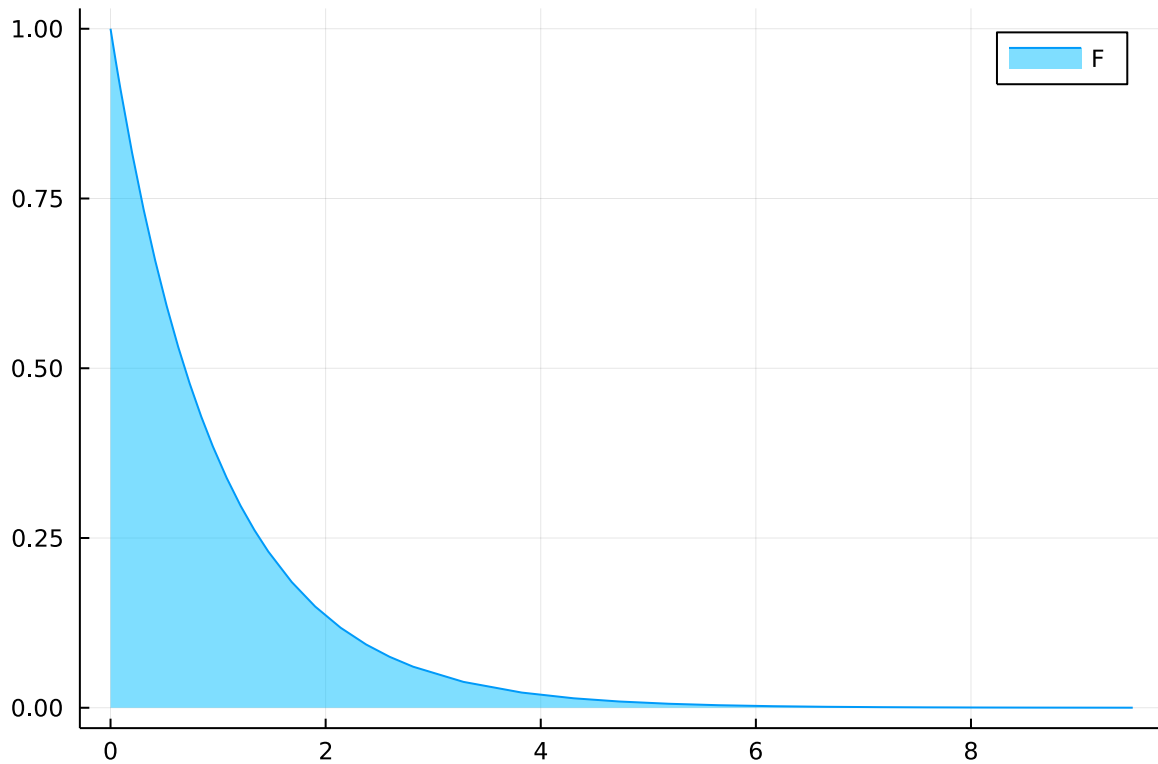
F distribution

2

```
• @bind  $\alpha$  Slider(2:5, show_value = true)
```

295

```
• @bind  $\beta$  Slider(295:305, show_value = true)
```



```
• StatsPlots.plot(Distributions.FDist( $\alpha$ ,  $\beta$ ), fill = (0, 0.5, :deepskyblue), label = "F")
```

Data

Generate random values for mass for three groups, I, II, III. Mass values will be taken from normal distribution. The total sample size will be 300, with 100 subjects in each group.

```
n = 300
```

```
• n = 300 # Sample size
```

```
n_group = 100
```

```
• n_group = Int(n / 3)
```

```
id = 1:300
```

```
• id = 1:n
```

```

• group = vcat(
•   repeat(["I"], n_group),
•   repeat(["II"], n_group),
•   repeat(["III"], n_group)
• );

```

10

```

• begin
•   μ1 = 98
•   μ2 = 100
•   μ3 = 105
•   σ1 = 20
•   σ2 = 15
•   σ3 = 10
• end

```

[104.777, 118.715, 68.505, 151.497, 80.9689, 90.469, 97.1297, 134.322, 87.616, 114.919, 1

```

• begin
•   Random.seed!(12) # Reproducible results
•   mass = vcat(
•     rand(Distributions.Normal(μ1, σ1), n_group),
•     rand(Distributions.Normal(μ2, σ2), n_group),
•     rand(Distributions.Normal(μ3, σ3), n_group)
•   )
• end

```

df =

	ID	Group	Mass
1	1	"I"	104.777
2	2	"I"	118.715
3	3	"I"	68.505
4	4	"I"	151.497
5	5	"I"	80.9689
6	6	"I"	90.469
7	7	"I"	97.1297
8	8	"I"	134.322
9	9	"I"	87.616
10	10	"I"	114.919
more			
300	300	"III"	94.6326

```

• df = DataFrames.DataFrame(
•   ID = id,
•   Group = group,
•   Mass = mass
• )

```

It may be useful to generate sub-DataFrame objects, one for each group. To do this, we use the filter function.

	ID	Group	Mass
1	201	"III"	110.947
2	202	"III"	120.308
3	203	"III"	116.931
4	204	"III"	100.125
5	205	"III"	111.328
6	206	"III"	116.786
7	207	"III"	111.071
8	208	"III"	92.9178
9	209	"III"	109.014
10	210	"III"	105.275
more			
100	300	"III"	94.6326

```
• begin
•   group_I = filter(r -> r.Group == "I", df)
•   group_II = filter(r -> r.Group == "II", df)
•   group_III = filter(r -> r.Group == "III", df)
• end
```

Although we generated the data manually, we extract the mass values from the DataFrame below as Vector objects.

```
• md"Although we generated the data manually, we extract the mass values from the
  DataFrame below as Vector objects."
```

```
[110.947, 120.308, 116.931, 100.125, 111.328, 116.786, 111.071, 92.9178, 109.014, 105.275
```

```
• begin
•   mass_all = collect(df.Mass) # All subjects
•   mass_I = collect(group_I.Mass) # Only group I subjects
•   mass_II = collect(group_II.Mass) # Only group II subjects
•   mass_III = collect(group_III.Mass) # Only group III subjects
• end
```

Summary statistics

```
mean_mass = 101.73033851002305
```

```
• mean_mass = Statistics.mean(mass_all)
```

```
mean_mass_I = 99.71311763141634
```

```
• mean_mass_I = mean(mass_I)
```

```
mean_mass_II = 100.52462781132863
```

```
• mean_mass_II = mean(mass_II)
```

```
mean_mass_III = 104.95327008732414
```

```
• mean_mass_III = mean(mass_III)
```

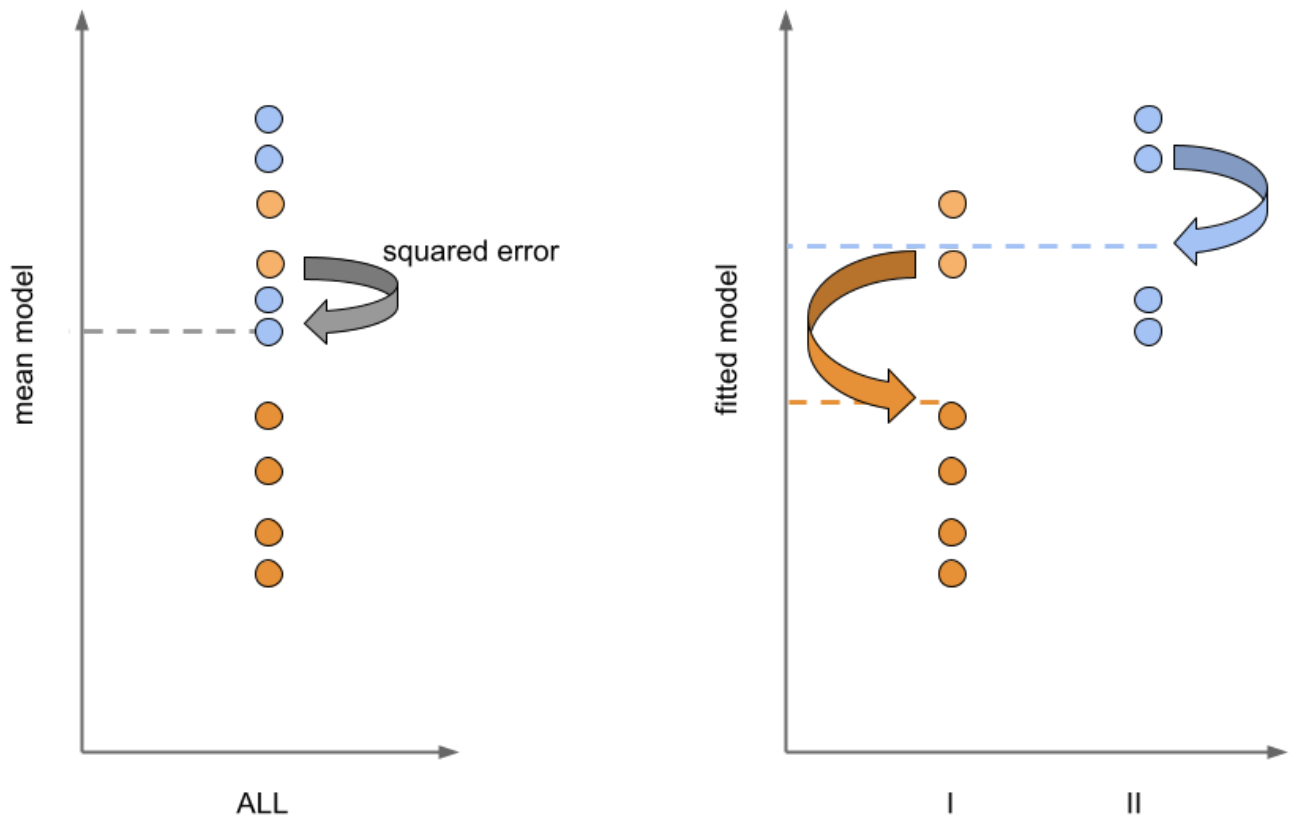
***F* ratio**

=====

To compare these three (or more) means, we need a statistic. We use the *F* statistic. The *F* statistic requires a few values and two parameters. The values and parameters stem from two models (and hence the *ratio* between them).

In the image below, we see two models. The first (on the left) has the values for a variable for all the subjects. The *mean* model uses the overall mean as prediction for the variable. If we square the difference between each value and the model prediction (the mean), we get a squared error. Summing over all these squared errors gives us the sum of squared errors for the *mean* (worst) model. We denote this as the SSM.

On the right, we see the variable separated by group. In the image there are only two groups. Our example has three. Irrespective of the number of groups, we can build a model, where the mean of each group is the predicted value for that group. We can then calculate a separate sum of squared errors for each and then sum over all the squared errors. This is termed the sum of squared errors for the fitted model. We denote this by SSF.



• `PlutoUI.LocalResource("SquaredErrors.png")`

As an example of the equation for the sum of squared errors we consider the SSM. We calculate a mean for the variable of interest for all the subjects combined. If we consider a vector of all these variable values \mathbf{y} then we express the mean as \bar{y} . The sum of squared errors for the mean (SSM) is the sum of the squared differences between the mean \bar{y} and each value of the variable y_i , where i is $1, 2, \dots, n$ and n is the sample size. This is shown in (1).

$$SSM = \sum_{i=1}^n (\bar{y} - y_i)^2 \quad (1)$$

The two parameters required for the F distribution relate to degrees of freedom. The equation for the F ratio has a numerator and denominator (each of which has a numerator and denominator). In the (overall) numerator, we have the difference between the number of parameters of the fitted model (3 in our case since we have three groups) and the number of parameters in the mean model (which is 1 since all subjects are placed together). In the denominator we have the difference between the overall sample size n and the number of parameters in the fitted model.

The equation for the F ratio is given below in (2).

$$F = \frac{\frac{SSM - SSF}{p_{\text{fit}} - p_{\text{mean}}}}{\frac{SSF}{n - p_{\text{fit}}}} \quad (2)$$

Below, we save the two parameter values.

3

```
• begin
•   p_mean = 1
•   p_fit = 3
• end
```

72297.54252117896

```
• begin
•   ssm = sum((mean_mass .- mass_all).^2)
•   ssf_I = sum((mean_mass_I .- mass_I).^2)
•   ssf_II = sum((mean_mass_II .- mass_II).^2)
•   ssf_III = sum((mean_mass_III .- mass_III).^2)
•   ssf = ssf_I + ssf_II + ssf_III
• end
```

Below we assign the F ratio for our data to the computer variable `f_ratio` using equation (2).

```
f_ratio = 3.2679750310137083
```

```
• f_ratio = ((ssm - ssf) / (p_fit - p_mean)) / ((ssf) / (n - p_fit))
```

Now we can express how likely this F ratio is.

Generating a sampling distribution by reassignment under the null hypothesis

.....

Under the null hypothesis, which states that all three means are equal, we can randomly reassign subjects to the three groups. At each reassignment, we recalculate the F ratio and append that to an (initially) empty vector, `f_ratio_sampling_5000`. This *builds* a sampling distribution of F ratio values. Below, we do this random reassignment 5000 times using a `for` loop.

```
reassign = 5000
```

```
• reassign = 5000 # Number of reassignments
```

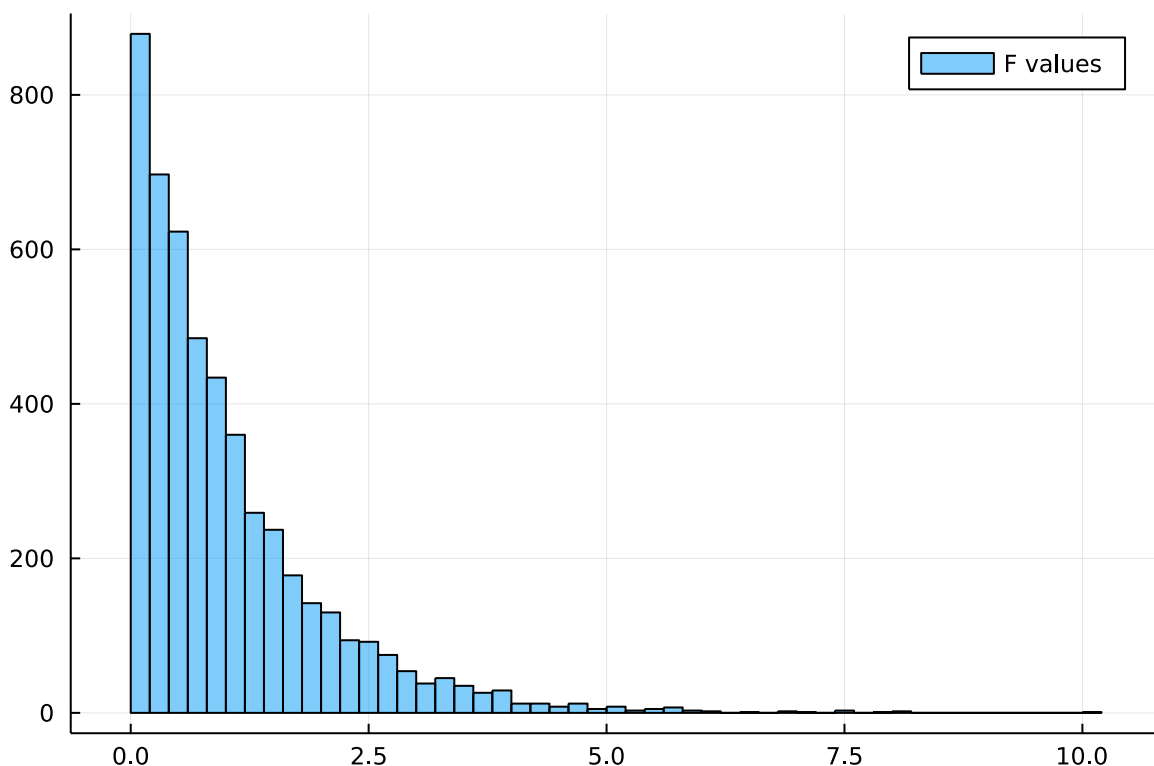
```
• f_ratio_sampling_5000 = []; # Empty vector to hold all 20000 values
```

```

• for i in 1:reassign
•     shuffle_mass = Random.shuffle(mass_all)
•     mean_shuffle_mass = mean(shuffle_mass)
•     ss_mean_reassign = sum((mean_shuffle_mass .- shuffle_mass).^2)
•
•     new_group_I = shuffle_mass[1:100]
•     mean_new_group_I = mean(new_group_I)
•     new_group_II = shuffle_mass[101:200]
•     mean_new_group_II = mean(new_group_II)
•     new_group_III = shuffle_mass[201:300]
•     mean_new_group_III = mean(new_group_III)
•
•     fit_I = sum((mean_new_group_I .- new_group_I).^2)
•     fit_II = sum((mean_new_group_II .- new_group_II).^2)
•     fit_III = sum((mean_new_group_III .- new_group_III).^2)
•     ss_fit_reassign = fit_I + fit_II + fit_III
•
•     reassign_f_ratio = ((ss_mean_reassign - ss_fit_reassign) / (p_fit - p_mean)) /
•     ((ss_fit_reassign) / (n - p_fit))
•
•     append!(f_ratio_sampling_5000, reassign_f_ratio)
• end

```

We can create a histogram of all the F ratio's. Compare this to the initial plot, setting the parameters to 2 and 297.



```

• StatsPlots.histogram(f_ratio_sampling_5000, nbins = 60, fillalpha = 0.5, label = "F
  values")

```

Estimating probability of our F ratio

Finally, we can express how likely it was to have *found* our particular F ratio by expressing the fraction of resampled F ratio values were larger than ours.

0.041

```
• sum(f_ratio_sampling_5000 .> f_ratio) / reassign
```

There we have it!